

Part Two: Decoding: The Viterbi Algorithm

Zihao Zhu, Gangze Li, Yufeng Xie

Tsinghua University

November 5, 2019

Decoding

Decoding: Given as input an HMM $\lambda = (A, B)$ and a sequence of observations $O = o_1, o_2, \dots, o_T$, find the most probable sequence of states $Q = q_1 q_2 q_3 \dots q_T$.





Figure: In the ice-cream domain, given a sequence of ice-cream observations 3 1 3 and an HMM, the task of the decoder is to find the best hidden weather sequence (H H H).

Problem

- ▶ One possibility
 - ▶ For each hidden state sequence Q
 - ▶ HHH, HHC, HCH, etc.
 - ▶ Compute $P(O|Q)$
 - ▶ Pick the highest one
- ▶  likelihood: N^T
a set of N states, a sequence of T observations
- ▶ Instead
 - ▶ Viterbi algorithm

Problem

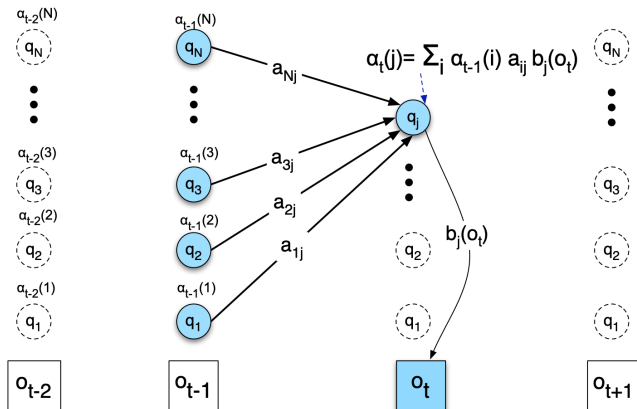


Figure: exponentially large number of state sequences! ¹

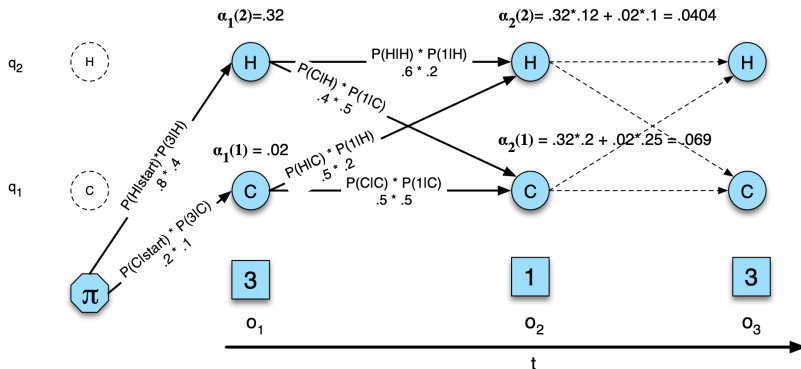
¹example from [JM19]

Viterbi algorithm

- ▶ The Viterbi algorithm is a dynamic programming algorithm for finding the most likely sequence of hidden states.
- ▶ We want to compute the joint probability of the observation sequence together with the best state sequence.
 - ▶ $v_t(j) = P(q_0, q_1, \dots, q_{t-1}, o_1, o_2, \dots, o_t, q_t = j | \lambda)$
 - ▶ $\alpha_t(j) = P(o_1, o_2 \dots o_t, q_t = j | \lambda)$

The forward trellis


compute the total observation likelihood for the ice-cream events 3 1 3



$$\alpha_t(j) = P(o_1, o_2 \dots o_t, q_t = j | \lambda)$$

Viterbi algorithm

the computation of $v_t(j)$ for two states at two time steps



$$v_t(j) = P(q_0, q_1, \dots, q_{t-1}, o_1, o_2, \dots, o_t, q_t = j | \lambda)$$

Viterbi algorithm

- ▶ The computation in each cell follows
$$v_t(j) = \max_{1 \leq i \leq N-1} v_{t-1}(i) a_{ij} b_j(o_t)$$
- ▶ The resulting probability expressed in each cell is
$$v_t(j) = P(q_0, q_1, \dots, q_{t-1}, o_1, o_2, \dots, o_t, q_t = j | \lambda)$$

the Viterbi backtrace

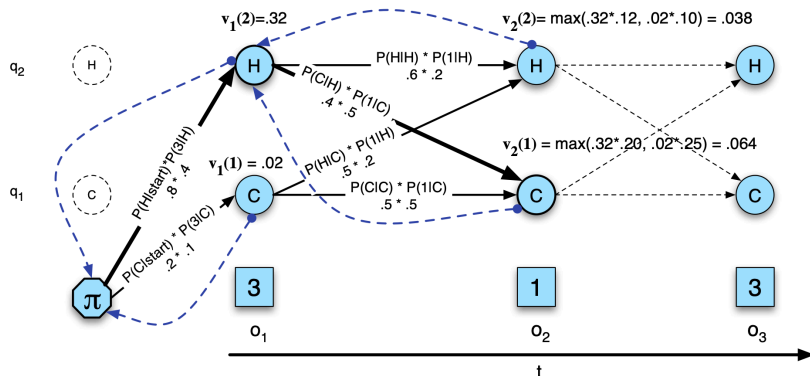


Figure: The Viterbi trellis for computing the best path through the hidden state space for the ice-cream eating events 3 1 3.

Viterbi recursion

► Initialization

$$\begin{aligned}v_1(j) &= \pi_j b_j(o_1) & 1 \leq j \leq N \\bt_1(j) &= 0 & 1 \leq j \leq N\end{aligned}$$

► Recursion

$$\begin{aligned}v_t(j) &= \max_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t); & 1 \leq j \leq N, 1 < t \leq T \\bt_t(j) &= \operatorname{argmax}_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t); & 1 \leq j \leq N, 1 < t \leq T\end{aligned}$$

► Termination

$$\begin{aligned}\text{The best score: } P_* &= \max_{i=1}^N v_T(i) \\ \text{The start of backtrace: } q_{T_*} &= \operatorname{argmax}_{i=1}^N v_T(i)\end{aligned}$$

Pseudocode

```
function VITERBI(observations of len  $T$ , state-graph of len  $N$ ) returns best-path, path-prob

create a path probability matrix viterbi[ $N, T$ ]
for each state  $s$  from 1 to  $N$  do                                ; initialization step
     $viterbi[s, 1] \leftarrow \pi_s * b_s(o_1)$ 
     $backpointer[s, 1] \leftarrow 0$ 
for each time step  $t$  from 2 to  $T$  do                            ; recursion step
    for each state  $s$  from 1 to  $N$  do
         $viterbi[s, t] \leftarrow \max_{s'=1}^N viterbi[s', t-1] * a_{s', s} * b_s(o_t)$ 
         $backpointer[s, t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s', t-1] * a_{s', s} * b_s(o_t)$ 

 $bestpathprob \leftarrow \max_{s=1}^N viterbi[s, T]$                         ; termination step
 $bestpathpointer \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s, T]$             ; termination step
 $bestpath \leftarrow$  the path starting at state  $bestpathpointer$ , that follows  $backpointer[]$  to states back in time
return  $bestpath$ ,  $bestpathprob$ 
```

Three fundamental problems [Rab89]

- ▶ **Problem 1 (Likelihood)** Given an HMM $= (A, B)$ and an observation sequence O , determine the likelihood $P(O|\lambda)$.
- ▶ **Problem 2 (Decoding)** Given an observation sequence O and an HMM $= (A, B)$, discover the best hidden state sequence Q .
- ▶ **Problem 3 (Learning)** Given an observation sequence O and the set of states in the HMM, learn the HMM parameters A and B .

Application²



The GENSCAN Web Server at MIT

Identification of complete gene structures in genomic DNA



[For information about Genscan, click here](#)

Server update, November, 2009: We've been recently upgrading the GENSCAN webserver hardware, which resulted in some problems in the output of GENSCAN. We apologize for the inconvenience. These output errors were resolved.

This server provides access to the program Genscan for predicting the locations and exon-intron structures of genes in genomic sequences from a variety of organisms.

This server can accept sequences up to 1 million base pairs (1 Mbp) in length. If you have trouble with the web server or if you have a large number of sequences to process, request a local copy of the program (see instructions at the bottom of this page).

Organism: Suboptimal exon cutoff (optional):

Sequence name (optional):

Print options:

Upload your DNA sequence file (upper or lower case, spaces/numbers ignored): no file selected

Or paste your DNA sequence here (upper or lower case, spaces/numbers ignored):

²from the link

GENSCAN

GENSCAN was developed by Chris Burge in the research group of Samuel Karlin, Department of Mathematics, Stanford University [BK97, SSK98].

Burge and Karlin in their gene finding computer program GENSCAN used a three-periodic (inhomogeneous) fifth order Markov chain to model coding regions of DNA sequences [BE06].

CpG island³

CATTCCGCTTCTCTCCCGAGGTGGCGCGTGGGA
GGTGTTTTGGCTCGGGTTCTGTAAGAATAGGCCAGG
CAGCTTCCCGCGGGATGCGCTCATCCCTCTCGG
GGTTCCGCTCCACCGCGCGCTTCGCGCGGTT
CGGCTGCGAGATGTTTTCCGACGGACAATGATTC
CACTCTCGCGCCCTCCCATGTTGATCCAGCTCCT
CTGCGGGCGTCAGGACCCCTGGGCCCGCCCCCG
CTCCACTCAGTCAATCTTTGTCCCGTATAAGGCG
GATTATCGGGGTGGCTGGGGGCGGCTGATTCGGA
CGAATGCCCTTGGGGGTACCCCGGGAGGGAACCT
CGGGCTCGCTTTGGCCAGCCCGCACCCCTGGT
TGAGCGGGCGAGAGGCCACAGGGGGCGCTCG
ATGTTCTGCAGCCCCCGCAGCAGCCCCACTCC
CGGCTCACCTCAAGATTGGCTGGCCCGCCCGAG
CTCTGTGCTGTGATTGGTCACAGCCCGTGTCCGT
CGCGCGCGCGGCGGATACAGGTGACCGCGCA
GAGGCCAGCTCGGGCGGTGTCCCGCGCGCGG
GACTCGGGCGGAGTTTCGCGAGGGCGGAAGCG
GGGAGTGTGCGGAGCGGTCTGGGAGGCGC
CGCGCGCGCTCGGAGCAGCTCCCCTCTCCGCA
GCCGTCACCGCGCGCGTCCCGCGCCCTGGCC
TCCCGCACTCGCGCACTCCTGTCCCGCGCCACG
CGCCACCTCCCACTCGATCGGTGCGGGCTGC
TGCGTGTGGGGCTCGGAGCGCGCCCTGCGG
CTCGCGCGCGCTGCTCGCGCTGAGGTGCGT
CGGTGCCCGGCCCGCGCGCCCGCGCGCGCG
GGCTCTGTTGACCGGTGCGCGCGTGTCTGC
AGCGGGGCTGAGGTAAGCGCGGGGCTGGCG
CGGTTGGCGCGCGGTGCGCGGGTTGGGAGGG
GGCGCGCTTCGCGCGGGAGGAGCGCGCGGCCG
GGTCGGGCGGGGTCTGAGGGGA


CTCTTAGTTTTGGGTGCATTTGTCTGGTCTTCCAA
CTAGATTGAAAGCTCTGAAAAAATATCTTGT
GTTTCTATCTGTTGAGCTCATAGTATCCAGGA
AGTAGTAGGGTTGACTGCATTGATTGGGACTACAC
TGGGAGTTTTCTTCCCATCTCCCTTAGTTTTCT
TTTTTCTTTCTTTCTTTCTTTTTTCTTTTTTT
TTGAGATGTCTGCTCAGTCCCCAGGCTGGA
GTGCAGTGGTGCGATCTTGGCTCACTGTAGCTCC
ACCTCCAGGTTCAAGCAATCTACTGCCTTAGCCT
CCCGAGTAGCTGGGATTACAAGCACCGCCACCAT
TCCTGGCTAAATTTTTTTTGTATTTTGTAGTGAGA
CAGGGTTTACCATGTTGGTGTGCTGGTCTCAGA
CTCCTGGGGCTAGCGATCCCCCTGCTCAGCCT
CCCAGAGTGTAGGATTACAGGCATGAGCCACTGT
ACCCGGCTCTCTCCAGTTTCCAGTTGGAATCCAA
GGGAAGTAAGTTTAAGATAAAGTTACGATTTTGA
CTTTGGATTGAGAAGAAATTTGTACCTTTAACACCT
AGAGTTGAACTTCATACCTGGAGAGCCTTAACATT
AAGCCCTAGCCAGCCTCCAGCAAGTGGCAATGGT
CAGGTTTGGCAGGATTCTCCCTGAAGTGGACT
GAGAGCCACACCCCTGGCCTGCACCATCCATCC
CCTATCCTTAGTGAAGCAAACTCCTTTGTTCCCTT
CTCCTTCTCCTAGTGACAGGAAATATTGTATCCTA
AAGAATGAAATAGCTTTGTACACCTCGTGGCCTCAG
GCCTCTTGACTTCAGGCGTCTGTATTAATCAAGT
GACATCTCCCGAGGCTCCCTGAATGTGGCAGATG
AAAGAGACTAGTTCAACCTGACCTGAGGGGAAAG
CCTTTGTGAAGGGTCAGGAG

Left: CpG sites at 1/10 nucleotides, constituting a CpG island. The sample is of a gene-promoter, the highlighted ATG consitutes the start codon.

Right: CpG sites present at every 1/100 nucleotides, constituting a more normal example of the genome, or a region of the genome that is commonly methylated.

³from the link

Toy Example ⁴




This model is composed of 2 states, H (high GC content) and L (low GC content). We can for example consider that state H characterizes coding DNA while L characterizes non-coding DNA.

$$p_H(A, 4) = e_H(A) \max(p_L(C, 3)p_{LH}, p_H(C, 3)p_{HH})$$


⁴example from the link

Toy Example



We used here $\log_2(p)$, and compute sums instead of products.


Toy Example



GGCA**T**GAA

- ▶ Probability (in \log_2) that G at the first position was emitted by state H, $p_H(G, 1) = -1 - 1.737 = -2.737$
- ▶ Probability (in \log_2) that G at the first position was emitted by state L, $p_L(G, 1) = -1 - 2.322 = -3.322$


Toy Example



GGCACTGAA

- Probability (in \log_2) that G at the 2nd position was emitted by state H,
 $p_H(G, 2) = -1.737 + \max(p_H(G, 1) + p_{HH}, p_L(G, 1) + p_{LH})$
 $= -1.737 + \max(-2.737 - 1, -3.322 - 1.322)$
 $= -5.474$ (obtained from $p_H(G, 1)$)
- Probability (in \log_2) that G at the 2nd position was emitted by state L,
 $p_L(G, 2) = -2.322 + \max(p_H(G, 1) + p_{HL}, p_L(G, 1) + p_{LL})$
 $= -2.322 + \max(-2.737 - 1.322, -3.322 - 0.737)$
 $= -6.059$ (obtained from $p_H(G, 1)$)

Toy Example




GGCACTGAA

	G	G	C	A	C	T	G	A	A
H	-2.73	-5.47	-8.21	-11.53	-14.01	...			-25.65
L	-3.32	-6.06	-8.79	-10.94	-14.01	...	→	→	→24.49

We then compute iteratively the probabilities $p_H(i, x)$ and $p_L(i, x)$ that nucleotide i at position x was emitted by state H or L , respectively. The highest probability obtained for the nucleotide at the last position is the probability of the most probable path. This path can be retrieved by back-tracking.

Toy Example



GGCACTGAA






back-tracking

(= finding the path which corresponds to the highest probability, -24.49)

	G	G	C	A	C	T	G	A	A
H	-2.73	-5.47	-8.21	-11.53	-14.01	...			-25.65
L	-3.32	-6.06	-8.79	-10.94	-14.01	...	→	→	-24.49

The most probable path is: **HHHLLLLL**

References I

-  Mark Borodovsky and Svetlana Ekiheva, *Problems and solutions in biological sequence analysis*, Cambridge University Press, 2006.
-  Chris Burge and Samuel Karlin, *Prediction of complete gene structures in human genomic dna*, Journal of molecular biology **268** (1997), no. 1, 78–94.
-  Dan Jurafsky and James H. Martin, *Speech and language processing (3rd ed. draft)*, 2019.
-  Lawrence R Rabiner, *A tutorial on hidden markov models and selected applications in speech recognition*, Proceedings of the IEEE **77** (1989), no. 2, 257–286.
-  SL Salzberg, DB Searls, and S Kasif, *Modeling dependencies in pre-mrna splicing signals*, Computational methods in molecular biology **32** (1998), 129.

References II